

Plan Bouquets: Query Processing without Selectivity Estimation

Anshuman Dutt

Database Systems Lab, SERC/CSA

Indian Institute of Science, Bangalore, INDIA

anshuman@dsl.serc.iisc.ernet.in

Jayant R. Haritsa

Database Systems Lab, SERC/CSA

Indian Institute of Science, Bangalore, INDIA

haritsa@dsl.serc.iisc.ernet.in

ABSTRACT

Selectivity estimates for optimizing OLAP queries often differ significantly from those actually encountered during query execution, leading to poor plan choices and inflated response times. We propose here a conceptually new approach to address this problem, wherein the compile-time estimation process is completely eschewed for error-prone selectivities. Instead, a small “bouquet” of plans is identified from the set of optimal plans in the query’s selectivity error space, such that at least one among this subset is near-optimal at each location in the space. Then, at run time, the actual selectivities of the query are incrementally “discovered” through a sequence of partial executions of bouquet plans, eventually identifying the appropriate bouquet plan to execute. The duration and switching of the partial executions is controlled by a graded progression of isocost surfaces projected onto the optimal performance profile. We prove that this construction results in bounded overheads for the selectivity discovery process and consequently, guaranteed worst-case performance. In addition, it provides repeatable execution strategies across different invocations of a query.

The plan bouquet approach has been empirically evaluated on both PostgreSQL and a commercial DBMS, over the TPC-H and TPC-DS benchmark environments. Our experimental results indicate that, even with conservative assumptions, it delivers substantial improvements in the worst-case behavior, without impairing the average-case performance, as compared to the native optimizers of these systems. Moreover, the bouquet technique can be largely implemented using existing optimizer infrastructure, making it relatively easy to incorporate in current database engines.

Overall, the bouquet approach provides novel guarantees that open up new possibilities for robust query processing.

Categories and Subject Descriptors

H.2.4 [Database Management Systems]: Query Processing

Keywords

Selectivity Estimation; Plan Bouquets; Robust Query Processing

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SIGMOD’14, June 22–27, 2014, Snowbird, UT, USA.

Copyright 2014 ACM 978-1-4503-2376-5/14/06 ...\$15.00.

<http://dx.doi.org/10.1145/2588555.2588566>.

1. INTRODUCTION

Cost-based database query optimizers estimate a host of *selectivities* while identifying the ideal execution plan for declarative OLAP queries. For example, consider **EQ**, the simple SPJ query shown in Figure 1 for enumerating orders of cheap parts – here, the optimizer estimates the selectivities of a selection predicate ($p_{\text{retailprice}}$) and two join predicates ($\text{part} \bowtie \text{lineitem}$, $\text{lineitem} \bowtie \text{orders}$). In practice, these estimates are often significantly in error with respect to the actual values subsequently encountered during query execution. Such errors, which can even be in *orders of magnitude* in real database environments [18], arise due to a variety of well-documented reasons [23], including outdated statistics, attribute-value independence(AVI) assumptions, coarse summaries, complex user-defined predicates, and error propagation in the query execution operator tree [16]. Moreover, in environments such as ETL workflows, the statistics may actually be *unavailable* due to data source constraints, forcing the optimizer to resort to “magic numbers” for the values (e.g. 1/10 for equality selections [22]). The net result of these erroneous estimates is that the execution plans recommended by the query optimizer may turn out to be poor choices at run-time, resulting in substantially inflated query response times.

```
select * from lineitem, orders, part
where p_partkey = l_partkey and l_orderkey =
o_orderkey and p_retailprice < 1000
```

Figure 1: Example Query (EQ)

A considerable body of literature exists on proposals to tackle this classical problem. For instance, techniques for improving the *statistical quality* of the meta-data include improved summary structures [1, 19], feedback-based adjustments [23], and on-the-fly re-optimization of queries [17, 4, 20]. A complementary approach is to identify *robust plans* that are relatively less sensitive to estimation errors [9, 3, 4, 14]. While these prior techniques provide novel and innovative formulations, they are limited in their scope and performance, as explained later in the related work section.

Plan Bouquet Approach

In this paper, we investigate a conceptually new approach, wherein the compile-time estimation process is completely eschewed for error-prone selectivities. Instead, these selectivities are systematically *discovered* at run-time through a calibrated sequence of cost-limited plan executions. In a nutshell, we attempt to side-step the selectivity estimation problem, rather than address it head-on, by adopting a “*seeing is believing*” viewpoint on these values.

1D Example. We introduce the new approach through a restricted 1D version of the EQ example query wherein only the `p_retailprice` selection predicate is error-prone. First, through repeated invocations of the optimizer, we identify the “parametric optimal set of plans” (POSP) that cover the entire selectivity range of the predicate. A sample outcome of this process is shown in Figure 2, wherein the POSP set is comprised of plans P1 through P5. Further, each plan is annotated with the selectivity range over which it is optimal – for instance, plan P3 is optimal in the (1.0%, 7.5%] interval. (In Figure 2, P = Part, L = Lineitem, O = Order, NL = Nested Loops Join, MJ = Sort Merge Join, and HJ = Hash Join).

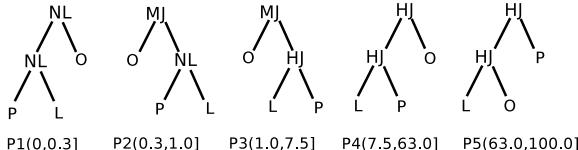


Figure 2: POSP plans on `p_retailprice` dimension

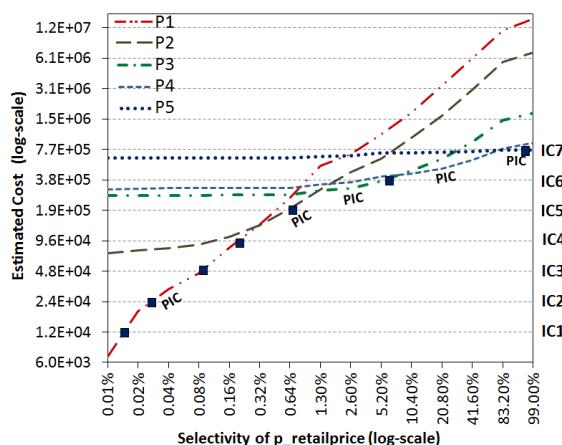


Figure 3: POSP performance (log-log scale)

The optimizer-generated costs of these POSP plans over the selectivity range are shown (on a log-log scale) in Figure 3. On this figure, we first construct the “POSP infimum curve” (PIC), defined as the trajectory of the minimum cost from among the POSP plans – this curve represents the ideal performance. The next step, which is a distinctive feature of our approach, is to *discretize* the PIC by projecting a graded progression of *isocost* (IC) steps onto the curve. For example, in Figure 3, the dotted horizontal lines represent a geometric progression of isocost steps, IC1 through IC7, with each step being *double* the preceding value. The intersection of each IC with the PIC (indicated by ■) provides an associated selectivity, along with the identity of the best POSP plan for this selectivity. For example, in Figure 3, the intersection of IC5 with the PIC corresponds to a selectivity of 0.65% with associated POSP plan P2. We term the subset of POSP plans that are associated with the intersections as the “plan bouquet” for the given query – in Figure 3, the bouquet consists of {P1, P2, P3, P5}.

The above exercises are carried out at query compilation time. Subsequently, at run-time, the correct query selectivities are explicitly discovered through a sequence of *cost-limited* executions of bouquet plans. Specifically, beginning with the cheapest isocost step, we iteratively execute the bouquet plan assigned to each step until either:

1. The partial execution overheads exceed the step’s cost value – in this case, we know that the actual selectivity location lies beyond the current step, motivating a switch to the next step in the sequence; or
2. The current plan completes execution within the budget – in this case, we know that the actual selectivity location has been reached, and a plan that is at least 2-competitive wrt the ideal choice was used for the final execution.

Example. To make the above process concrete, consider the case where the selectivity of `p_retailprice` is 5%. Here, we begin by partially executing plan P1 until the execution overheads reach IC1 (1.2E4 | 0.015%). Then, we extend our cost horizon to IC2, and continue executing P1 until the overheads reach IC2 (2.4E4 | 0.03%), and so on until the overheads reach IC4 (9.6E4 | 0.2%). At this juncture, there is a change of plan to P2 as we look ahead to IC5 (1.9E5 | 0.65%), and during this switching all the intermediate results (if any) produced thus far by plan P1 are *jettisoned*. The new plan P2 is executed till the associated overhead limit (1.9E5) is reached. The cost horizon is now extended to IC6 (3.8E5 | 6.5%), in the process jettisoning plan P2’s intermediate results and executing plan P3 instead. In this case, the execution will complete before the cost limit is reached since the actual location, 5%, is less than the selectivity limit of IC6. Viewed in toto, the net sub-optimality turns out to be 1.78 since the exploratory overheads are 0.78 times the optimal cost, and the optimal plan itself was (coincidentally) used for the final execution.

Extension to Multiple Dimensions. When the above approach is generalized to the multi-dimensional selectivity environment, the IC steps and the PIC curve become surfaces, and their intersections represent selectivity surfaces on which multiple bouquet plans may be present. For example, in the 2-D case, the IC steps are horizontal planes cutting through a hollow 3D PIC surface, typically resulting in hyperbolic intersection contours with different plans associated with disjoint segments of this contour – an instance of this scenario is shown in Figure 6.

Notwithstanding these changes, the basic mechanics of the bouquet algorithm remain virtually identical. The primary difference is that we jump from one IC surface to the next only after it is determined (either explicitly or implicitly) that *none* of the bouquet plans present on the current IC surface can completely execute the given query within the associated cost budget.

Performance Characteristics

At first glance, the plan bouquet approach, as described above, may appear to be utterly absurd and self-defeating because: (a) At compile time, considerable preprocessing may be required to identify the POSP plan set and the associated PIC; and (b) At run-time, the overheads may be hugely expensive since there are multiple plan executions for a single query – in the worst scenario, as many plans as are present in the bouquet!

However, we will attempt to make the case in the remainder of this paper, that it is indeed possible, through careful design, to have *plan bouquets efficiently provide robustness profiles that are markedly superior to the native optimizer’s profile*. Specifically, if we define robustness to be the worst-case sub-optimality in plan performance that can occur due to selectivity errors, the bouquet mechanism delivers substantial robustness improvements, while providing comparable or improved average-case performance.

For instance, the runtime performance of the bouquet technique on EQ is profiled in Figure 4 (dark blue curve). We observe that its

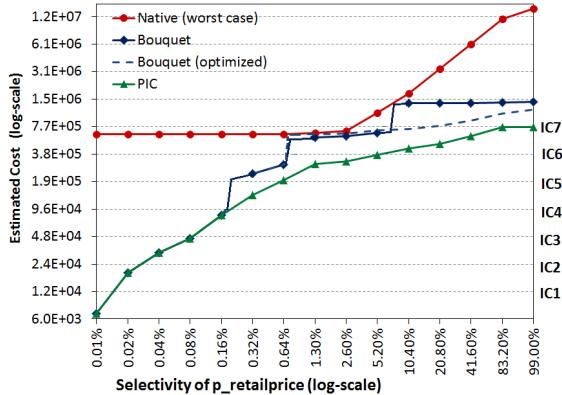


Figure 4: Bouquet Performance (log-log scale)

performance is much closer to the PIC (dark green) as compared to the worst case profile for the native optimizer (dark red), comprised of the supremum of the individual plan profiles. In fact, the worst case sub-optimality for the bouquet is only 3.6 (at 6.5%), whereas the native optimizer suffers a sub-optimality of around 100 when P5 (which is optimal for large selectivities) is mistakenly chosen to execute a query with a low selectivity of 0.01%. The average sub-optimality of the bouquet, computed over all possible errors, is 2.4, somewhat worse than the 1.8 obtained with the native optimizer. However, when the enhancements described later in this paper are incorporated, the optimized bouquet’s performance (dashed blue) improves to 3.1 (worst case) and 1.7 (average case), thereby dominating the native optimizer on both metrics.

Our motivation for the cost-based discretization of the PIC is that it leads to *guaranteed* bounds on worst-case performance. For instance, we prove that the cost-doubling strategy used in the 1D example results in an *upper-bound* of 4 for the worst-case sub-optimality – this bound is inclusive of all exploratory overheads incurred by the partial executions, and is irrespective of the query’s actual selectivity. In fact, we can go further to show that 4 is the best competitive factor achievable by *any* deterministic algorithm. For the multi-dimensional case, the bound becomes 4 times the bouquet cardinality (more accurately, the plan cardinality of the densest contour), and we present techniques to limit this cardinality to a small value. To our knowledge, these robustness bounds are the first such guarantees to be presented in the database literature (although similar characterizations are well established in the algorithms community [8]). Further, we also present a variety of design optimizations that result in a practical performance which is well within the theoretical bounds.

In order to empirically validate its utility, we have evaluated the bouquet approach on PostgreSQL and a popular commercial DBMS. Our experiments utilize a rich set of complex decision support queries sourced from the TPC-H and TPC-DS benchmarks. The query workload includes selectivity spaces with as many as five error-prone dimensions, thereby capturing environments that are extremely challenging from a robustness perspective. Our performance results indicate that the bouquet approach typically provides *orders of magnitude* improvements, as compared to the optimizer’s native choices. As a case in point, for Query 19 of the TPC-DS benchmark with 5 error prone join selectivities, the worst-case sub-optimality plummeted from about 10^6 to just 10! The potency of the approach is also indicated by the fact that for many queries, the bouquet’s average performance is within 4 times of the corresponding PICs.

What is even more gratifying is that the above performance profiles are *conservative* since we assume that at every plan switch, *all* previous intermediate results are completely thrown away – in practice, it is conceivable that some of these prior results could be retained and reused in the execution of a future plan.

Apart from improving robustness, there is another major benefit of the bouquet mechanism: On a given database, the execution strategy for a particular query instance, i.e. the sequence of plan executions, is *repeatable* across different invocations of the query instance – this is in marked contrast to prior approaches wherein plan choices are influenced by the current state of the database statistics and the query construction. Such **stability of performance is especially important for industrial applications, where considerable value is attributed to reproducible performance characteristics [3]**.

Finally, with regard to implementation, the bouquet technique can be largely constructed using techniques (e.g. abstract plan costing) that have already found expression in modern DB engines, as explained later in Section 5.4.

Thus far, we had tacitly assumed the optimizer’s *cost model* to be perfect – that is, only *optimizer costs* were used in the evaluations. While this assumption is certainly not valid in practice, improving the model quality is, in principle, an orthogonal problem to that of estimation. Notwithstanding, we also analyze the robustness guarantees in the presence of bounded modeling errors. Moreover, to positively verify robustness improvements, explicit run-time evaluations are also included in our experimental study.

In closing, we wish to highlight that from a deployment perspective, the bouquet technique is intended to *complementarily co-exist* with the classical optimizer setup, leaving it to the user or DBA to make the choice of which system to use for a specific query instance – **essential factors that are likely to influence this choice are discussed in the epilogue.**

Organization. The remainder of the paper is organized as follows: In Section 2, a precise description of the robust execution problem is provided, along with the associated notations. Theoretical bounds on the robustness provided by the bouquet technique are presented in Section 3. We then discuss its design methodology, the compile-time aspects in Section 4 and the run-time mechanisms in Section 5. The experimental framework and performance results are reported in Section 6. Related work is reviewed in Section 7, while Section 8 presents a critical review of the bouquet approach.

2. PROBLEM FRAMEWORK

In this section, we present our robustness model, the associated performance metrics, and the notations used in the sequel. Robustness can be defined in many different ways and there is no universally accepted metric [13] – here, we use the notion of performance sub-optimality to characterize robustness.

The error-prone selectivity space is denoted as ESS and its dimensionality D is determined by the number of error-prone selectivity predicates in the query. The space is represented by a grid of D-dimensional points with each point $q(s_1, s_2, \dots, s_D)$ corresponding to a unique query with selectivity s_j on the j^{th} dimension. The cost of a plan P_i at a query location q in ESS is denoted by $c_i(q)$.

For simplicity, we assume that the estimated query locations and the actual query locations are uniformly and independently distributed over the entire discretized selectivity space – that is, all estimates and errors are equally likely. This definition can easily be extended to the general case where the estimated and actual locations have idiosyncratic probability distributions.

Given a user query Q , denote the optimizer’s *estimated* location of this query by q_e and the *actual* location at runtime by q_a . Next,

denote the plan chosen by the optimizer at q_e as P_{oe} , and the optimal plan at q_a by P_{oa} . With these definitions, the sub-optimality incurred due to using P_{oe} at q_a is simply defined as the ratio:

$$SubOpt(q_e, q_a) = \frac{c_{oe}(q_a)}{c_{oa}(q_a)} \quad \forall q_e, q_a \in ESS \quad (1)$$

with $SubOpt$ ranging over $[1, \infty)$. The worst-case $SubOpt$ for a given query location q_a is defined to be wrt the q_e that results in the maximum sub-optimality, that is, where selectivity inaccuracies have the maximum adverse performance impact:

$$SubOpt_{worst}(q_a) = \max_{q_e \in ESS} (SubOpt(q_e, q_a)) \quad \forall q_a \in ESS \quad (2)$$

With the above, the global worst-case is simply defined as the (q_e, q_a) combination that results in the maximum value of $SubOpt$ over the entire ESS, that is,

$$MSO = \max_{q_a \in ESS} (SubOpt_{worst}(q_a)) \quad (3)$$

Further, given the uniformity assumption about the distribution of estimated and actual locations, the average sub-optimality over ESS is defined as:

$$ASO = \frac{\sum_{q_e \in ESS} \sum_{q_a \in ESS} SubOpt(q_e, q_a)}{\sum_{q_e \in ESS} \sum_{q_a \in ESS} 1} \quad (4)$$

The above MSO and ASO definitions are appropriate for the way that modern optimizers behave, wherein selectivity estimates are made at compile-time, and a single plan is executed at run-time. However, in the plan bouquet technique, neither of these characteristics is true – error-prone selectivities are not estimated at compile-time, and multiple plans may be invoked at run-time. Notwithstanding, we can still compute the corresponding statistics by: (a) substituting q_e with a “don’t care” *; (b) replacing P_{oe} with P_b to denote the plan bouquet mechanism; and (c) having the cost of the bouquet, $c_b(q_a)$, include the overheads incurred by the exploratory partial executions. Further, the running selectivity location, as progressively discovered by the bouquet mechanism, is denoted by q_{run} .

Even when the bouquet algorithm performs well on the MSO and ASO metrics, it is possible that for some specific locations $q_a \in ESS$, it performs poorer than the worst performance of the native optimizer – it is therefore harmful for the queries associated with these locations. This possibility is captured using the following *MaxHarm* metric:

$$MH = \max_{q_a \in ESS} \left(\frac{SubOpt(*, q_a)}{SubOpt_{worst}(q_a)} - 1 \right) \quad (5)$$

Note that MH values lie in the range $(-1, MSO_{bouquet} - 1]$ and harm occurs whenever MH is positive.

An assumption that fundamentally underlies the entire bouquet mechanism is that of *Plan Cost Monotonicity* (PCM) – that is, the costs of the POSP plans increase monotonically with increasing selectivity values. This assumption has often been made in the literature [5, 6, 15], and holds for virtually all the plans generated by PostgreSQL on the benchmark queries. The only exception we have found is for queries featuring *existential* operators, where the POSP plans may exhibit *decreasing* monotonicity with selectivity. Even in such scenarios, the basic bouquet technique can be utilized by the simple expedient of plotting the ESS with $(1 - s)$ instead of s on the selectivity axes. Thus, only queries whose optimal cost surfaces have a maxima or minima in the *interior* of the error space, are not amenable to our approach.

3. ROBUSTNESS BOUNDS

We begin our presentation of the plan bouquet approach by characterizing its performance bounds with regard to the MSO metric, initially for the 1D scenario, and then extending it to the general multi-dimensional case.

3.1 1D Selectivity Space

As described in the Introduction, the 1D PIC curve is discretized by projecting a graded progression of isocost steps onto the curve. We assume that the PIC is an increasing function (by virtue of PCM) and continuous throughout ESS; its minimum and maximum costs are denoted by C_{min} and C_{max} , respectively. Now, specifically consider the case wherein the isocost steps are organized in a *geometric* progression with initial value a ($a > 0$) and common ratio r ($r > 1$), such that the PIC is sliced with $m = \log_r \lceil \frac{C_{max}}{C_{min}} \rceil$ cuts, IC_1, IC_2, \dots, IC_m , satisfying the boundary conditions $a/r < C_{min} \leq IC_1$ and $IC_{m-1} < C_{max} = IC_m$, as shown in Figure 5.

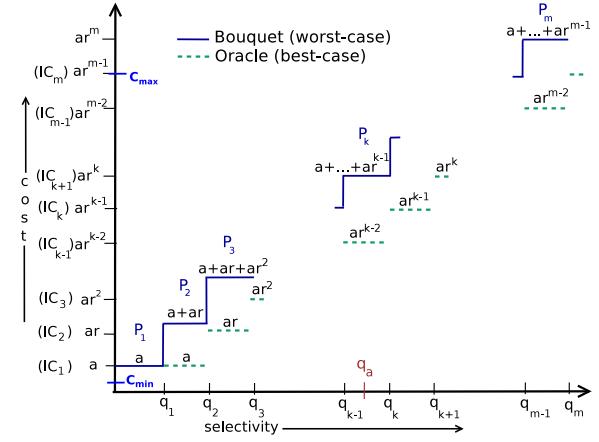


Figure 5: 1D Selectivity Space

For $1 \leq k \leq m$, denote the selectivity location where the k^{th} isocost step (IC_k) intersects the PIC by q_k and the corresponding bouquet plan as P_k . All the q_k locations are unique by definition due to the PCM and continuity requirements on the PIC curve. However, it is possible that some of the P_k plans may be common to multiple intersection points (e.g. in Figure 3, plan P_1 was common to steps IC_1 through IC_4). Finally, for mathematical convenience, assign q_0 to be 0.

With this framework, the bouquet execution algorithm operates as follows in the most general case, where a different plan is associated with each step: We start with plan P_1 and budget IC_1 , progressively working our way up through the successive bouquet plans P_2, P_3, \dots until we reach the first plan P_k that is able to fully execute the query within its assigned budget IC_k . It is easy to see that the following lemma holds:

LEMMA 1. If q_a resides in the range $(q_{k-1}, q_k]$, $1 \leq k \leq m$, then plan P_k executes it to completion in the bouquet algorithm.

PROOF. We prove by contradiction: If q_a was located in the region $(q_k, q_{k+1}]$, then P_k could not have completed the query due to the PCM restriction. Conversely, if q_a was located in $(q_{k-2}, q_{k-1}]$, P_{k-1} itself would have successfully executed the query to completion. With similar reasoning, we can prove the same for the remaining regions that are beyond q_{k+1} or before q_{k-2} .

Performance Bounds. Consider the generic case where q_a lies in the range $(q_{k-1}, q_k]$. Based on Lemma 1, the associated worst case cost of the bouquet execution algorithm is given by the following expression:

$$\begin{aligned} C_{\text{bouquet}}(q_a) &= \text{cost}(IC_1) + \text{cost}(IC_2) + \dots + \text{cost}(IC_k) \\ &= a + ar + ar^2 + \dots + ar^{k-1} = \frac{a(r^k - 1)}{r - 1} \end{aligned} \quad (6)$$

The corresponding cost for an ‘oracle’ algorithm that magically apriori knows the correct location of q_a is lower bounded by ar^{k-2} , due to the PCM restriction. Therefore, we have

$$\text{SubOpt}(*, q_a) \leq \frac{\frac{a(r^k - 1)}{r - 1}}{\frac{r^2}{ar^{k-2}}} = \frac{r^2}{r - 1} - \frac{r^{2-k}}{r - 1} \leq \frac{r^2}{r - 1} \quad (7)$$

Note that the above expression is *independent* of k , and hence of the specific location of q_a . Therefore, we can state for the entire selectivity space, that:

THEOREM 1. *Given a query Q on a 1D error-prone selectivity space, and the associated PIC discretized with a geometric progression having common ratio r, the bouquet execution algorithm ensures that: $\text{MSO} \leq \frac{r^2}{r - 1}$*

Further, the choice of r can be optimized to minimize this value – the RHS reaches its minima at $r = 2$, at which the value of MSO is 4. The following theorem shows that this is the *best* performance achievable by any deterministic online algorithm – leading us to conclude that the *doubling* based discretization is the ideal solution.

THEOREM 2. *No deterministic online algorithm can provide an MSO guarantee lower than 4 in the 1D scenario.*

PROOF. We prove by contradiction, assuming there exists an optimal online robust algorithm, R^* with a MSO of f , $f < 4$.

Firstly, note that R^* must have a monotonically increasing sequence of plan execution costs, $a_1, a_2, \dots, a_{k^*+1}$ in its quest to find a plan P_{k^*+1} that can execute the query to completion. The proof is simple: If $a_i > a_j$ with $i < j$, then we could construct another algorithm that skips the a_j execution and still execute the query to completion using P_{k^*+1} , and therefore has less cumulative overheads than R^* , which is not possible by definition.

Secondly, if R^* stops at P_{k^*+1} , then q_a has to necessarily lie in the range $(q_{k^*}, q_{k^*+1}]$ (Lemma 1 holds for any monotonic algorithm). Therefore, the worst-case performance of R^* is given by $\sum_{i=1}^{k^*+1} a_i \leq f$. Since q_a could be chosen to lie in any interval, this inequality should hold true across all intervals, i.e.

$$\forall j \in 1, 2, \dots, k^*: \frac{\sum_{i=1}^{j+1} a_i}{a_j} \leq f$$

Using the notation A_j to represent $\sum_{i=1}^j a_i$ and Y_j to represent the ratio $\frac{A_{j+1}}{A_j}$, we can rewrite the above as:

$$\frac{A_{j+1}}{a_j} \leq f \Rightarrow A_{j+1} \leq f(A_j - A_{j-1}) \Rightarrow \frac{A_{j+1}}{A_j} \leq f \frac{(A_j - A_{j-1})}{A_j}$$

that is, $Y_j \leq f(1 - \frac{1}{Y_{j-1}})$.

We can show through elementary algebra that $\forall z > 0$, $(1 - \frac{1}{z}) \leq \frac{z}{4}$. Therefore, we have that $Y_j \leq (\frac{f}{4})Y_{j-1}$, leading to $Y_{k^*} \leq (\frac{f}{4})^{k^*-1} Y_1$. Using the assumption of $f < 4$, we can find a sufficiently large k^* such that $(\frac{f}{4})^{k^*-1} Y_1 < 1$. Hence, $Y_{k^*} < 1$ which implies that $A_{k^*+1} < A_{k^*}$, a contradiction.

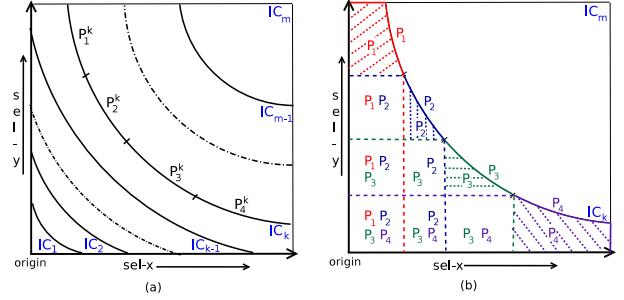


Figure 6: 2D Selectivity Space

3.2 Multi-dimensional Selectivity Space

We now move to the general case of multi-dimensional selectivity error spaces. A sample 2D scenario is shown in Figure 6a, wherein the isocost surfaces IC_k are represented by *contours* that represent a continuous sequence of selectivity locations (in contrast to the single location in the 1D case). Further, multiple bouquet plans may be present on each individual contour as shown for IC_k wherein four plans, $P_1^k, P_2^k, P_3^k, P_4^k$, are the optimizer’s choices over disjoint x, y selectivity ranges on the contour. Now, to decide whether q_a lies below or beyond IC_k , in principle *every* plan on the IC_k contour has to be executed – only if none complete, do we know that the actual location definitely lies beyond the contour.

This need for exhaustive execution is highlighted in Figure 6b, where for the four plans lying on IC_k , the regions in the selectivity space on which each of these plans is guaranteed to complete within the IC_k budget are enumerated (the contour superscripts are omitted in the figure for visual clarity). Note that while several regions are ‘covered’ by multiple plans, each plan also has a region that it alone covers – the hashed regions in Figure 6b. For queries located in such regions, only the execution of the associated unique plan would result in confirming that the query is within the contour.

The basic bouquet algorithm for the multi-dimensional case is shown in Figure 7, using the notation n_k to represent the number of plans on contour k .

```

for cid = 1 to m do                                ▷ for each cost-contour cid
    for i = 1 to ncid do                          ▷ for each plan on cid
        start executing Picid
        while running-cost(Picid) ≤ cost-budget(ICcid) do
            execute plan Picid                      ▷ cost limited execution
            if Picid finishes execution then
                return query result
            stop executing Picid

```

Figure 7: Multi-dimensional Bouquet Algorithm

Performance Bounds. Given a query Q with q_a located in the range $(IC_{k-1}, IC_k]$, the worst-case total execution cost for the multi-D bouquet algorithm is given by

$$C_{\text{bouquet}}(q_a) = \sum_{i=1}^k [n_i \times \text{cost}(IC_i)] \quad (8)$$

Using ρ to denote the number of plans on the *densest* contour, and upper-bounding the values of the n_i with ρ , we get the following performance guarantee:

$$C_{\text{bouquet}}(q_a) \leq \rho \times \sum_{i=1}^k \text{cost}(IC_i) \quad (9)$$

Now, following a similar derivation as for the 1D case, we arrive at the following theorem:

THEOREM 3. *Given a query Q with a multidimensional error-prone selectivity space, the associated PIC discretized with a geometric progression having common ratio r and maximum contour plan density ρ , the bouquet execution algorithm ensures that:*

$$MSO \leq \rho \frac{r^2}{r-1}$$

Setting $r = 2$ in this expression ensures that $MSO \leq 4\rho$.

3.3 Minimizing IsoCost Surface Plan Density

To the best of our knowledge, the above MSO bounds are the first such guarantees in the literature. While the 1D bounds are inherently strong giving a guarantee of 4 or better, the multi-dimensional bounds, however, depend on ρ , the maximum plan density over the isocost surfaces. Therefore, to have a practically useful bound, we need to ensure that the value of ρ is kept to the minimum.

This can be achieved through the *anorexic reduction* technique described in [15]. Here, POSP plans are allowed to “swallow” other plans, that is, occupy their regions in the ESS space, if the sub-optimality introduced due to these swallowings can be bounded to a user-defined threshold, λ . In [15], it was shown that even for complex OLAP queries, a λ value of 20% was typically sufficient to bring the number of POSP plans down to “anorexic levels”, that is, a small absolute number within or around 10.

When we introduce the anorexic notion into the bouquet setup, it has two opposing impacts on the sub-optimality guarantees – on the one hand, the constant multiplication factor is increased by a factor $(1 + \lambda)$; on the other, the value of ρ is significantly reduced. Overall, the deterministic guarantee is altered from $4\rho_{POSP}$ to $4(1 + \lambda)\rho_{ANOREXIC}$.

Empirical evidence that this tradeoff is very beneficial is shown in Table 1, which compares for a variety of multi-dimensional error spaces, the bounds (using Equation 8) under the original POSP configuration and under an anorexic reduction ($\lambda = 20\%$). As a particularly compelling example, consider 5D_DS_Q19, a five-dimensional selectivity error space based on Q19 of TPC-DS – we observe here that the bound plummets by more than an order of magnitude, going down from 379 to 30.4.

Error Space	ρ_{POSP}	MSO Bound	$\rho_{ANOREXIC}$	MSO Bound
3D_H_Q5	11	33	3	12.0
3D_H_Q7	13	34	3	9.6
4D_H_Q8	88	213	7	24.0
5D_H_Q7	111	342.5	9	37.2
3D_DS_Q15	7	23.5	3	12.0
3D_DS_Q96	6	22.5	3	13.0
4D_DS_Q7	29	83	4	17.8
4D_DS_Q26	25	76	5	19.8
4D_DS_Q91	94	240	9	35.3
5D_DS_Q19	159	379	8	30.4

Table 1: Performance Guarantees (POSP versus Anorexic)

3.4 Cost Modeling Errors

Thus far, we had catered to arbitrary errors in selectivity estimation, but assumed that the cost model itself was perfect. In practice, this is certainly not the case, but if the modeling errors were to be unbounded, it appears hard to ensure robustness since, in principle, the estimated cost of any plan could be arbitrarily different to the actual cost encountered at run-time. However, we could think of an intermediate situation wherein the modeling errors are non-zero but *bounded* – specifically, the estimated cost of any plan, given

correct selectivity inputs, is known to be within a δ error factor of the actual cost. That is, $\frac{c_{estimated}}{c_{actual}} \in [\frac{1}{(1 + \delta)}, (1 + \delta)]$.

Our construction is lent credence to by the recent work of [24], wherein static cost model tuning was explored in the context of PostgreSQL – they were able to achieve an average δ value of around 0.4 for the TPC-H suite of queries.

This “unbounded estimation errors, bounded modeling errors” framework is amenable to robustness analysis – specifically, it is easy to show that (proof in [12])

$$MSO_{bounded_modeling_error} \leq MSO_{perfect_model} * (1 + \delta)^2$$

So, for instance, when $\delta = 0.4$, corresponding to the average in [24], the MSO increases by at most a factor of 2.

4. BOUQUET: COMPILE-TIME

In this section, we describe the compile-time aspects of the bouquet algorithm, whose complete work-flow is shown in Figure 8.

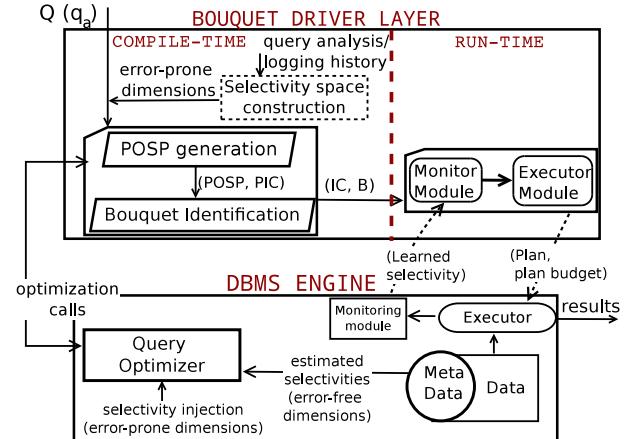


Figure 8: Architecture of Bouquet Mechanism

4.1 Selectivity Space Construction

Given a user query Q, the first step is to identify the error-prone selectivity dimensions in the query. For this purpose, we can leverage the approach proposed in [17], wherein a set of uncertainty modeling rules are outlined to classify selectivity errors into categories ranging from “no uncertainty” to “very high uncertainty”. Alternatively, a log could be maintained of the errors encountered by similar queries in the workload history. Finally, there is always the fallback option of making *all* predicates where selectivities are evaluated, to be selectivity dimensions for the query.

The chosen dimensions form the ESS selectivity space. In general, each dimension ranges over the entire [0,100] percentage range – however, due to schematic constraints, the range may be reduced. For instance, the maximum legal value for a PK-FK join is the reciprocal of the PK relation’s minimum row cardinality.

4.2 POSP Generation

The next step is to determine the parametric optimal set of plans (POSP) over the entire ESS. Producing the complete POSP set requires repeated invocations of the query optimizer at a high degree of resolution over the space. This process can, in principle, be computationally very expensive, especially for higher-dimensional spaces. However, user queries are often submitted

through “canned” form-based interfaces – for such environments it appears feasible to offline *precompute* the entire POSP set.

Further, even when this is not the case, the overheads can be made manageable by leveraging the following observation: The full POSP set is not required, only the subset that lies on the isocost surfaces. Therefore, we begin by optimizing the two locations at the corners of the principal diagonal of the selectivity space, giving us C_{min} and C_{max} . From these values, the costs of all the isocost contours are computed. Then, the ESS is divided into smaller hypercubes, recursively dividing those hypercubes through which one or more isocost contours pass – a contour passes through a hypercube if its cost is within the cost range established by the corners of the hypercube’s principal diagonal. The recursion stops when we reach hypercubes whose sizes are small enough that it is cheap to explicitly optimize all points within them. In essence, only a narrow “band” of locations around each contour is optimized.

Finally, note that the POSP generation process is “embarrassingly parallel” since each location in the ESS can be optimized independent of the others. Therefore, hardware resources in the form of multi-processor multi-core platforms can also be leveraged to bring the overheads down to practical levels.

Selectivity Injection. As discussed above, we need to be able to systematically generate queries with the desired ESS selectivities. One option is to, for each new location, suitably modify the query constants and the data distributions, but this is clearly impractically cumbersome and time-consuming. We have therefore taken an alternative approach in our PostgreSQL implementation, wherein the optimizer is instrumented to directly support *injection* of selectivity values in the cost model computations. Interestingly, some commercial optimizer APIs already support such selectivity injections to a limited extent (e.g. IBM DB2 [26]).

4.3 Plan Bouquet Identification

Armed with knowledge of the plans on each of the isocost contour surfaces, which is usually in the several tens or hundreds of plans, the next step is to carry out a cost-based anorexic reduction [15] in order to bring the plan cardinality down to a manageable number. That is, we identify a smaller set of plans, such that each replaced location now has a new plan whose cost is within $(1+\lambda)$ times the optimal cost. We denote the set of plans on the surface of IC_k with B_k and the union of these sets of plans provides the final plan bouquet i.e. $B = \cup_{k=1}^m B_k$. Finally, the isocost surfaces (IC), annotated with their updated costs (the original costs are inflated by $1 + \lambda$ to account for the anorexic reduction), and B , the set of bouquet plans, are passed to the run-time phase.

5. BOUQUET: RUN TIME

In this section, we present the run-time aspects of the bouquet mechanism, as per the work-flow shown in Figure 8.

The basic bouquet algorithm (Figure 7) discovers the location of a query by sequentially executing the set of plans on each contour in a cost-limited manner until either one of them completes, or the plan set is exhausted, forcing a jump to the next contour. Note that in this process, no explicit monitoring of selectivities is required since the execution statuses serve as implicit indicators of whether we have reached q_a or not. However, as we will show next, consciously tracking selectivities can aid in substantively curtailing the discovery overheads. In particular, the tracking can help to (a) reduce the number of plan executions incurred in crossing contours; and (b) develop techniques for increasing the selectivity movement obtained through each cost-limited plan execution.

5.1 Reducing Contour Crossing Executions

In this optimization, during the processing of a contour, the location of q_{run} is *incrementally* updated after each (partial) plan execution to reflect the additional knowledge gained through the execution. An example learning sequence is shown in Figure 9 – here, the q_{run} known at the conclusion of IC_{k-1} is progressively updated via q_{run}^1 and q_{run}^2 to reach q_{run}^3 on IC_k , with the corresponding plan execution sequence being P_1, P_4, P_3 (the contour superscripts are omitted for ease of exposition). The important point to observe here is that the contour crossing was accomplished *without executing P_2* .

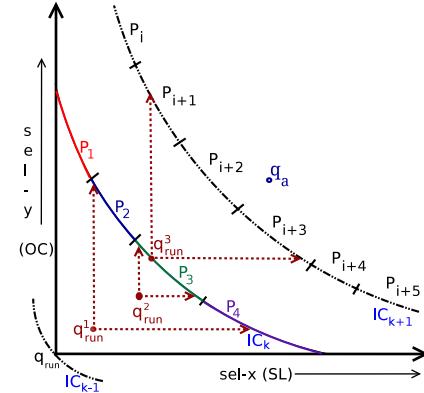


Figure 9: Minimizing Contour Crossing Executions

We now discuss how the plan execution sequence is decided. The strategy used is to ensure that at all times, the actual location is in the *first quadrant* with respect to the current location as origin – this invariant allows us to use the positive axes as a “pincer” movement towards reaching the desired target, in the process eliminating from consideration some plans on the contour. Specifically, at each q_{run} location, we first identify $AxisPlans$, the set of bouquet plans present at the *intersection* of the isocost contour with the dimensional-axes corresponding to q_{run} as origin. For example, in Figure 9, $AxisPlans(q_{run})$ is comprised of P_4 and P_1 , corresponding to the x and y dimensions, respectively. Then, from within this set, we heuristically pick, using a combination of structure and cost considerations, the plan that promises to provide the maximum movement towards q_a . The specific heuristic used is the following: The plans in $AxisPlans$ are first ordered based on their costs at q_{run} , and then clustered into “equivalence groups” based on the closeness of these costs. From the cheapest equivalence group, the plan with an error-prone node occurring deepest in the plan-tree is chosen for execution. The expectation is that being cheapest at q_{run} provides the maximum spare budget, while having error-prone nodes deep within the plan-tree ensures that this spare budget is not uselessly spent on processing error-free nodes.

In Figure 9, the above heuristic happens to choose P_1 at q_{run} and thereby reach q_{run}^1 . The process is repeated with q_{run} set to q_{run}^1 – now $AxisPlans(q_{run}^1)$ is $\{P_2, P_4\}$, and P_4 is chosen by the heuristic, resulting in a movement to q_{run}^2 . Finally, with q_{run} set to q_{run}^2 , $AxisPlans(q_{run}^2)$ contains only P_3 which is executed to reach q_{run}^3 , and hence IC_k . Note, as mentioned before, that P_2 is eliminated from consideration in this incremental process.

There is a further advantage of the incremental updates to q_{run} : When we reach q_{run}^3 in Figure 9, we not only learn that q_a lies beyond IC_k but can also ab initio eliminate 3 plans (P_i, P_{i+4}, P_{i+5}) from the list of candidate plans for crossing IC_{k+1} , since these plans lie outside the first quadrant of q_{run}^3 .

5.2 Monitoring Selectivity Movement

Having established the utility of incremental updates to q_{run} , we now go into the details of its implementation. Consider the scenario wherein Figure 9 represents the selectivity updation process for a TPC-H based query with error-prone join selectivities s_{SL} and s_{OC} on the x and y dimensions, respectively. Correspondingly, let plans P_1 through P_4 be as shown in Figure 10. Further, each node j of these plans is labeled with the corresponding tuple count, t_j , obtained at the end of the cost-limited execution – these annotations are explicitly shown for P_1 in Figure 10.

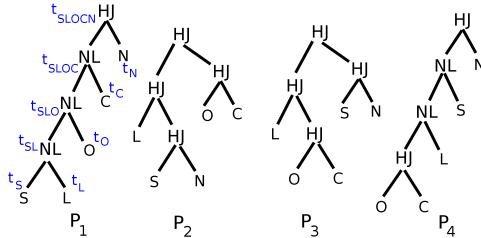


Figure 10: Plans on the k^{th} contour

After P_1 's execution, the tuple count on node SL can be utilized to update the running selectivity \hat{s}_{SL} as $\frac{t_{SL}}{|S|_e \times |L|_e}$ where $|S|_e$ and $|L|_e$ denote the cardinalities of the input relations to the SL join. The values in the denominator are clearly known before execution as these nodes are assumed to be error-free. Note that \hat{s}_{SL} is a *lower bound* on s_{SL} , and therefore continues to maintain the “first quadrant” invariant required by the bouquet approach.

The other selectivity s_{OC} , is not present as an independent node in plan P_1 . If we directly use $\hat{s}_{OC} = \frac{t_{SLOC}}{t_{SLO} \times |C|_e}$, there is a danger of *overestimation* wrt $q_a(OC)$ since t_{SLO} may not be known completely due to the cost-budgeted execution of P_1 . Such overestimations may lead to violation of the “first quadrant” property, and are therefore impermissible. Consequently, we defer the updating of s_{OC} to the subsequent execution of plans P_4 and P_3 where it can be independently computed from fully known inputs.

In general, given any plan-tree, we can learn the lower bound for an error-prone selectivity only after the cardinalities of its inputs are completely known. This is possible when either the inputs are apriori error-free, or any error-prone inputs have been completely learnt through the earlier executions. The latter method of learning allows the bouquet approach to function even in the (unlikely) case where it does not possess plans with independent appearances for all the error-prone selectivities. In the above discussion, an implicit assumption is that all selectivities are *independent* with respect to each other – this is in conformance with the typical modeling framework of current optimizers.

5.3 Maximizing Selectivity Movement

Now we discuss how individual cost-limited plan executions can be modified to yield maximum movement of q_{run} towards q_a in return for the overheads incurred in their partial executions – that is, to “get the maximum selectivity bang for the execution buck”.

In executing a budgeted plan to determine error-prone selectivities, we would ideally like the cost budget to be utilized as far as possible by the nodes in the plan operator tree that can provide us *useful* learning. However, there are two hurdles that come in the way: Firstly, the costs incurred by upstream nodes that *precede* the error nodes in the plan evaluation. Secondly, the costs incurred by the downstream nodes in the *pipeline* featuring the error nodes.

The first problem of upstream nodes can be ameliorated by preferentially choosing during the *AxisPlans* routine, as mentioned earlier, plans that feature the error-prone nodes deeper (i.e. earlier) in the plan-tree. The second problem of downstream nodes can be solved by deliberately breaking the pipeline immediately after first error node and *spilling* its output, which ensures that the downstream nodes do not get any data/tuples to process. These changes help to maximize the effort spent on executing the error-prone nodes, and thereby increase the selectivity movement with a given cost budget.

Movement Example. We now illustrate, using the same example scenario as Figure 9, as to how spill-based execution is utilized to achieve increased selectivity movement. In Figure 11, the spilled versions of the plans P_1 through P_4 are shown, denoted using \tilde{P} . The modified selectivity discovery process using the spilled partial executions is shown in Figure 12, with the progressive selectivity locations being q_{run}^a , q_{run}^b , q_{run}^c and q_{run}^d .

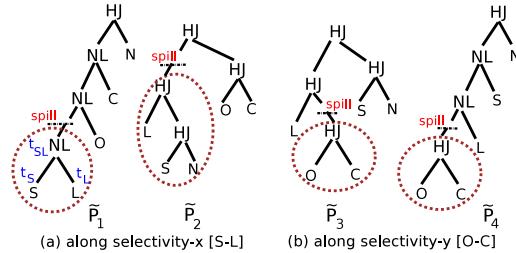


Figure 11: Plans (spilled version) and their movement direction

The discovery process starts with executing plan \tilde{P}_1 until its cost-limit is reached. The tuple count on the error-prone node SL is then used to calculate \hat{s}_{SL} , as discussed earlier. Since the budget allotted for the full plan is now solely focused on learning s_{SL} , it is reasonable to expect that there will be materially more movement in s_{SL} as compared to executing generic P_1 . In fact, it is easy to prove that, at the minimum, crossing of q_{run} from the third quadrant of the P_1 segment to its fourth quadrant is *guaranteed* – this minimal case is shown in Figure 12 as location q_{run}^a .

After \tilde{P}_1 exhausts its cost-budget, the *AxisPlans* routine chooses \tilde{P}_4 to take over, which starts learning s_{OC} , and ends up reaching at least q_{run}^b in Figure 12. Continuing in similar vein, \tilde{P}_2 is executed to reach q_{run}^c , and finally, \tilde{P}_3 is executed to reach q_{run}^d on the next contour. Due to focusing our energies on learning only a single selectivity in each plan execution, the movement of q_{run} follows a *Manhattan* profile from the origin upto q_a , as shown in Figure 12.

A high-level pseudocode of the full bouquet algorithm, incorporating the above optimizations, is presented in Figure 13.

5.4 Implementation Details

For implementing the bouquet mechanism, the database engine needs to support the following functionalities: (1) abstract plan costing; (2) selectivity injection during query optimization; (3) cost-limited partial execution of plans (generic and spilled); and (4) selectivity monitoring on a running basis. Abstract plan costing is supported by quite a few commercial engines including SQL Server [25], while limited selectivity injection is provided in DB2 [26]. The other two features were found to be easy to implement since they leverage pre-existing engine resources. For example, in PostgreSQL, the node-granularity tuple counter required for cost-limited execution, as well as selectivity monitoring, is available through the *instrumentation* data structure [29].

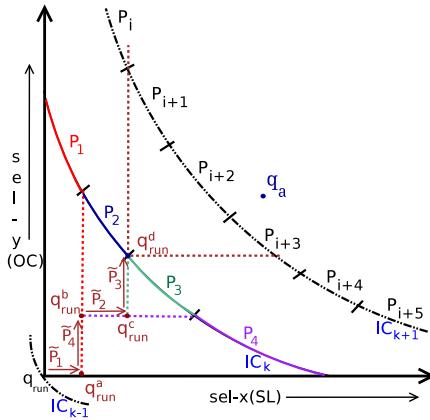


Figure 12: Maximizing Selectivity Movement

```

 $q_{run} = (0,0, \dots, 0); \quad cid = 1 \quad \triangleright initialization$ 
loop
   $P_{cur} = \text{AxisPlanRoutine}(q_{run}, cid) \quad \triangleright next plan selection$ 
  while running-cost( $P_{cur}$ )  $\leq$  cost-budget( $IC_{cid}$ ) do
    execute  $P_{cur} \quad \triangleright cost limited execution$ 
    if  $P_{cur}$  finishes execution then
      return query result
    update  $q_{run} \quad \triangleright selectivity updation$ 
    if optimal-cost( $q_{run}$ )  $\geq$  cost-budget( $IC_{cid}$ ) then
      cid ++  $\quad \triangleright early contour change$ 
  end while
end loop
  
```

Figure 13: Optimized Bouquet algorithm

5.5 Summary of Features

We complete this discussion of the mechanics of the bouquet approach with a synopsis of its distinctive features: (a) Compile-time estimation is completely eschewed for error-prone selectivities; (b) Plan switch decisions are triggered by predefined isocost contours (in contrast to dynamic criteria of [17, 18]); (c) Plan switch choices are restricted to an anorexic set of precomputed POSP plans; (d) AVI assumptions on intra-relational predicates are dispensed with since selectivities are explicitly monitored; (e) A first-quadrant invariant between the actual selectivity and the running selectivity is maintained, supporting monotonic progress towards the objective.

6. EXPERIMENTAL EVALUATION

We now turn our attention towards profiling the performance of the bouquet approach on a variety of complex OLAP queries, using the MSO, ASO and MH metrics enumerated in Section 2. As a precursor to these run-time metrics, we also discuss the overheads incurred by the bouquet algorithm in the compile-time phase.

Database Environment. The test queries (full descriptions in [12]) are chosen from the TPC-H and TPC-DS benchmarks to cover a spectrum of join-graph geometries, including *chain*, *star*, *branch*, etc. with the number of base relations ranging from 4 to 8. The number of error-prone selectivities range from 3 to 5 in these queries, all corresponding to join-selectivity errors, for making challenging multi-dimensional ESS spaces. We experiment with the TPC-H and TPC-DS databases at their default sizes of 1GB and 100GB, respectively, as well as larger scaled versions. Finally, the physical schema has indexes on all columns featuring in the queries, thereby maximizing the cost gradient $\frac{C_{max}}{C_{min}}$ and creating “hard-nut” environments for achieving robustness.

The summary query workload specifications are given in Table 2 – the naming nomenclature for the queries is xD_y_Qz , where x specifies the number of dimensions, y the benchmark (H or DS), and z the query number in the benchmark. So, for example, 3D_H_Q5 indicates a three-dimensional error selectivity space on Query 5 of the TPC-H benchmark.

Query	Join-graph (# relations)	$\frac{C_{max}}{C_{min}}$	Query	Join-graph (# relations)	$\frac{C_{max}}{C_{min}}$
3D_H_Q5	chain(6)	16	3D_DS_Q96	star(4)	185
3D_H_Q7	chain(6)	5	4D_DS_Q7	star(5)	283
4D_H_Q8	branch(8)	28	5D_DS_Q19	branch(6)	183
5D_H_Q7	chain(6)	50	4D_DS_Q26	star(5)	341
3D_DS_Q15	chain(4)	668	4D_DS_Q91	branch(7)	149

Table 2: Query workload specifications

System Environment. For the most part, the database engine used in our experiments is a modified version of PostgreSQL 8.4 [28], incorporating the changes outlined in Section 5.4. We also present sample results from a popular commercial optimizer. The hardware platform is a vanilla Sun Ultra 24 workstation with 8 GB memory and 1.2 TB of hard disk.

In the remainder of this section, we compare the bouquet algorithm (with anorexic parameter $\lambda = 20\%$) against the native PostgreSQL optimizer, and the SEER robust plan selection algorithm [14]. SEER uses a mathematical model of plan cost behavior in conjunction with anorexic reduction to provide replacement plans that, at all locations in ESS, either improve on the native optimizer’s performance, or are worse by at most the λ factor – it is therefore expected to perform better than the native optimizer on our metrics. It is important to note here that, in the SEER framework, the comparative yardstick is P_{oe} , the optimal plan at the *estimated* location, whereas in our work, the comparison is with P_{oa} , the optimal plan at the *actual* location.¹

For ease of exposition, we will hereafter refer to the bouquet algorithm, the native optimizer, and the SEER algorithm as **BOU**, **NAT** and **SEER**, respectively, in presenting the results.

6.1 Compile-time Overheads

The computationally expensive aspect of BOU’s compile-time phase is the identification of the POSP set of plans in ESS. For this task, we use the contour-focused approach described in Section 4, which ignores most of the space lying between contours. In all of our queries, the number of contours was no more than 10. Therefore, the contour-POSP was generated within a few hours even for 5D scenarios on our generic workstation, which appears a feasible investment for canned queries. Moreover, as described in Section 4.2, these overheads could be brought down to a few minutes, thanks to the inherent parallelism in the task.

6.2 Worst-case Performance (MSO)

In Figure 14, the MSO performance is profiled, on a log scale, for a set of 10 representative queries submitted to NAT, SEER and BOU. The first point to note is that NAT is *not* inherently robust – to the contrary, its MSO is huge, ranging from around 10^3 to 10^7 . Secondly, SEER also does not provide any material improvement on NAT – this may seem paradoxical at first glance, but is only to be expected once we realize that not *all* the highly sub-optimal (q_e, q_a) combinations in NAT were necessarily helped in

¹Purely heuristic-based reoptimization techniques, such as POP [18] and Rio [4], are not included in the evaluation suite since their performance could be arbitrarily poor with regard to both P_{oe} and P_{oa} , as explained in [12].

the SEER framework. Finally, and in marked contrast, BOU provides *orders of magnitude* improvements over NAT and SEER – as a case in point, for 5D_DS_Q19, BOU drives MSO down from 10^6 to around just 10. In fact, even in absolute terms, it consistently provides an MSO of *less than ten* across all the queries.

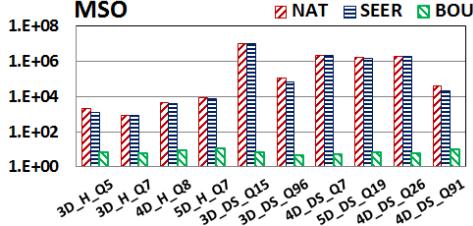


Figure 14: MSO Performance (log-scale)

6.3 Average-case Performance (ASO)

At first glance, it may be surmised that BOU’s dramatic improvement in worst-case behavior is purchased through a corresponding deterioration of average-case performance. To quantitatively demonstrate that this is not so, we evaluate ASO for NAT, SEER and BOU in Figure 15, again on a log scale. We see here that for some queries (e.g. 3D_DS_Q15), ASO of BOU is much better than that of NAT, while for the remainder (e.g. 4D_H_Q8) the performance is comparable. Even more gratifyingly, the ASO in absolute terms is typically less than 4 for BOU. On the other hand, SEER’s performance is again similar to that of NAT – this is an outcome of the high dimensionality of the error space which makes it extremely difficult to find universally safe replacements that are also substantively beneficial.

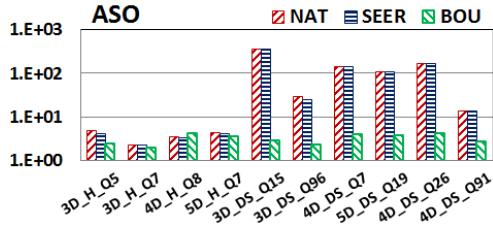


Figure 15: ASO Performance (log-scale)

6.4 Spatial Distribution of Robustness

We now profile for a sample query, namely 5D_DS_Q19, the percentage of locations for which BOU has a specific range of improvement over NAT. That is, the *spatial distribution* of enhanced robustness, $\frac{SubOpt_{worst}(q_a)}{SubOpt(*, q_a)}$. This statistic is shown in Figure 16, where we find that for the vast majority of locations (close to 90%), BOU provides *two or more orders of magnitude improvement* with respect to NAT. SEER, on the other hand, provides significant improvement over NAT for specific (q_e, q_a) combinations, but may not materially help the *worst-case* instance for each q_a . Therefore, we find that its robustness enhancement is less than 10 at all locations in the ESS.

6.5 Adverse Impact of Bouquet (MH)

Thus far, we have presented the improvements due to BOU. However, as highlighted in Section 2, there may be *individual* q_a

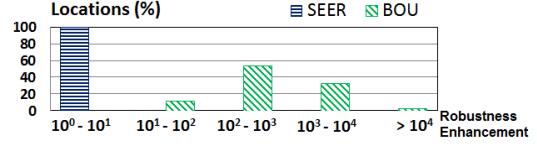


Figure 16: Distribution of enhanced Robustness (5D_DS_Q19)

locations where BOU performs poorer than NAT’s worst-case, i.e. $SubOpt(*, q_a) > SubOpt_{worst}(q_a)$. This aspect is quantified in Figure 17 where the maximum harm is shown (on a linear scale) for our query test suite. We observe that BOU may be up to a factor of 4 worse than NAT. Moreover, SEER steals a march over BOU since it *guarantees* that MH never exceeds λ (= 0.2). However, the important point to note is that the percentage of locations for which harm is incurred by BOU is less than 1% of the space. Therefore, from an overall perspective, the likelihood of BOU adversely impacting performance is rare, and even in these few cases the harm is limited (\leq MSO-1), especially when viewed against the order of magnitude improvements achieved in the beneficial scenarios.

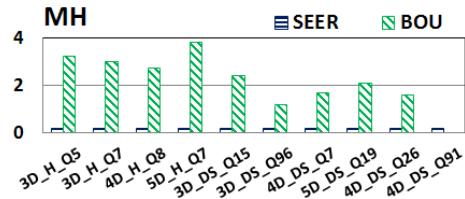


Figure 17: MaxHarm performance

6.6 Plan Cardinalities

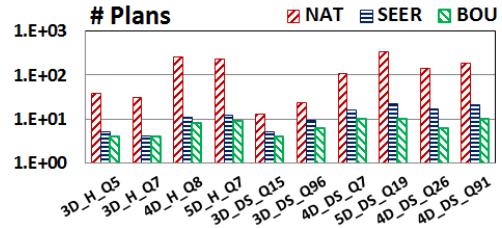


Figure 18: Plan Cardinalities (log-scale)

The plan cardinalities of NAT, SEER and BOU are shown on a log-scale in Figure 18. We observe here that although the original POSP cardinality may be in the several tens or hundreds, the number of plans in SEER is orders of magnitude lower, and those retained in BOU is even smaller – only around 10 or fewer, even for the 5D queries. This is primarily due to the initial anorexic reduction and the subsequent confinement to plan contours. The important implication of these statistics is that the bouquet size is, to the first degree of approximation, effectively *independent of the dimensionality and complexity of the error space*.

6.7 Query Execution Times (TPC-H)

To verify that the promised benefits of BOU are actually delivered at run-time, we also carried out experiments wherein query response times were explicitly measured for NAT and BOU. For this purpose, we crafted query instance 2D_H_Q8a (details in [12]), whose q_a was (33.7%, 45.6%), but NAT erroneously estimated the

location to be $q_e = (3.8\%, 0.02\%)$ due to incorrect AVI assumptions.² As a result, the plan chosen by NAT took almost 580 seconds to complete, whereas the optimal plan at q_a finished in merely 16 seconds, i.e $\text{SubOpt}(q_e, q_a) \approx 36$.

When BOU was invoked on the same 2D_H_Q8a query, it identified 6 bouquet plans spread across 7 isocost contours, resulting in an MSO bound of less than 20 (Equation 8). Subsequently, basic BOU produced the query result in about 117 seconds, involving 18 partial executions to cross 5 contours before the final full execution. Moreover, optimized BOU further brought the running time down to less than 70 seconds, using only 11 partial executions.

The isocost-contour-wise breakups of both basic and optimized BOU are given in Table 3, along with a comparative summary of their performance. Overall, the sub-optimality of optimized BOU is ≈ 4 , almost an order of magnitude lower than that of NAT (≈ 36). Note that the intended doubling of execution times across contours does not fully hold in Table 3 – this is an artifact of the imperfections in the underlying cost model of the PostgreSQL optimizer, compounded by our not having tuned this default model.

Contour ID	Avg Plan Exec. Time (in sec)	# Exec. (Basic BOU)	Time(sec) (Basic BOU)	# Exec. (Opt. BOU)	Time(sec) (Opt. BOU)
1	0.6	2	1.2	2	1.2
2	3.1	4	12.4	2	6.2
3	4.8	4	19.2	3	14.4
4	6.2	5	31.0	3	18.6
5	12.2	3	36.6	1	12.2
6	16.1	1	16.1	1	16.1
Total		19	116.5	12	68.7
Performance Summary (in seconds)		NAT	Basic BOU	Opt. BOU	Optimal
		579.4	116.5	68.7	16.1

Table 3: Bouquet execution for 2D_H_Q8a

6.8 Commercial Database Engine

All the results presented thus far were obtained on our instrumented PostgreSQL engine. We now present sample evaluations on a popular commercial engine, hereafter referred to as COM. Since COM’s API does not directly support injection of selectivities, we constructed queries 3D_H_Q5b and 4D_H_Q8b (details in [12]), wherein all error dimensions correspond to selection predicates on the base relations – the selectivities on such dimensions can be indirectly set up through changing only the constants in the query. The database and system environment remained identical to that of the PostgreSQL experiments.

Focusing on the performance aspects, shown in Figure 19, we find that here also large values of MSO and ASO are obtained for NAT and SEER. Further, BOU continues to provide substantial improvements on these metrics with a small sized bouquet. Again, the robustness enhancement is at least an order of magnitude for more than 90% of the query locations, without incurring any harm at the remaining locations ($MH < 0$). These results imply that our earlier observations are not artifacts of a specific engine.

7. RELATED WORK

A rich body of literature is available pertaining to selectivity estimation issues [11]. We start with the overview of the closely related techniques which can be collectively termed as *plan-switching approaches*, as they involve run-time switching among complete query plans. At first glance, our bouquet approach, with its partial

²We explicitly verified that there were no estimation errors in the remaining selectivity dimensions of the query.

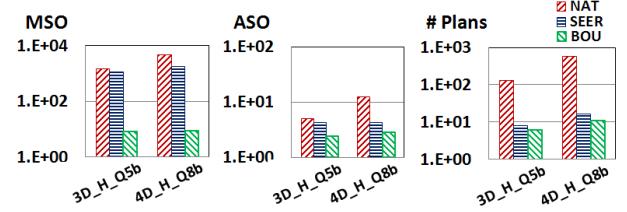


Figure 19: Commercial Engine Performance (log-scale)

execution of multiple plans, may appear very similar to run-time re-optimization techniques such as POP [18] and Rio [4]. However, there are key differences: Firstly, they start with the optimizer’s estimate as the initial seed, and then conduct a full-scale re-optimization if the estimate are found to be significantly in error. In contrast, we always start from the origin of the selectivity space, and directly choose plans from the bouquet for execution without invoking the optimizer again. A beneficial and unique side-effect of this start-from-origin approach is that it assures repeatability of the query execution strategy.

Secondly, both POP and Rio are based on heuristics and do not provide any performance bounds. In particular, POP may get stuck with a poor plan since its validity ranges are defined using structure-equivalent plans only. Similarly, Rio’s sampling-based heuristics for monitoring selectivities may not work well for join-selectivities and its definition of plan robustness on the basis of performance at corners (principal diagonal) has not been justified.

Recently, a novel interleaved optimization and execution approach was proposed in [20] wherein plan fragments are selectively executed, when recommended by an error propagation framework, to guard against the fallout of estimation errors. The error framework leverages an elegant histogram construction mechanism from [19] that minimizes the multiplicative error. While this technique substantively reduces the execution overheads, it provides no guarantees as it is largely based on heuristics.

Techniques that use a single plan during the entire query execution [9, 3, 14, 19, 6] run into the basic infeasibility of a single plan to be near-optimal across the entire selectivity space. The bouquet mechanism overcomes this problem by identifying a small set of plans that collectively provide the near-optimality property. Further, it does not require any prior knowledge of the query workload or the database contents. On the other hand, the use of only one active plan (at a time) to process the data makes the bouquet algorithm dissimilar from *Routing-based approaches* wherein different data segments may be routed to different simultaneously active plans – for example, plan per tuple [2] and plan per tuple group [21].

Our technique may superficially look similar to PQO techniques, (e.g. PPQO [5]), since a set of plans are identified before execution by exploring the selectivity space. The primary difference is that these techniques are useful for saving on optimization time for query instances with known parameters and selectivities. On the other hand, our goal is to regulate the worst case performance impact when the computed selectivities are likely to be erroneous.

Further, the bouquet technique does not modify plan structures at run-time (modulo spilling directives). This is a major difference from “plan-morphing” approaches, where the execution plan may be substantially modified at run-time using custom-designed operators, e.g. *chooseplan* [10], *switch* [4], *feedback* [7].

Finally, we emphasize that our goal of minimizing the worst case performance in the presence of unbounded selectivity errors, does not coincide with any of the earlier works in this area. Previously considered objectives include (a) improved performance compared to the optimizer generated plan [4, 14, 17, 18, 20]; (b) improved av-

verage performance and/or reduced variance [9, 6, 3]; (c) improved accuracy of selectivity estimation structures [1]; and (d) bounded impact of multiplicative estimation errors [19].

8. CRITIQUE OF BOUQUET APPROACH

Having presented the mechanics and performance of the bouquet approach, we now take a step back and critique the technique.

The bouquet approach is intended for use in difficult estimation environments – that is, in database setups where accurate selectivity estimation is hard to achieve. However, when estimation errors are apriori known to be small, re-optimization techniques such as [18, 4], which use the optimizer’s estimate as the initial seed, are likely to converge much quicker than the bouquet algorithm, which requires starting at the origin to ensure the first quadrant invariant. But, if the estimates were apriori guaranteed to be *under-estimates*, then the bouquet algorithm can also leverage the initial seed.

Being a *plan-switching* approach, the bouquet technique suffers from the drawbacks generic to such approaches: Firstly, they are poor at serving *latency-sensitive* applications as they have to enforce wait for the final plan execution to return result tuples. Secondly, they are not recommended for update queries since maintaining transactional consistency with multiple executions may incur significant overheads to rollback the effects of the aborted partial executions. Finally, with single-plan optimizers, DBAs use their domain knowledge to fine-tune the plan using “*plan-hints*”. But this is not straightforward in *plan-switching* techniques since the actual plan sequence is determined only at run-time. Notwithstanding the limitations, such techniques are now featured even in commercial products (e.g. [27]).

There are also a few problems that are *specific* to the bouquet approach: Firstly, while it is inherently robust to changes in data *distribution*, since these changes only shift the location of q_a in the existing ESS, the same is not true with regard to database *scale-up*. That is, if the database size increases significantly, then the original ESS no longer covers the entire error space. An obvious solution to handle this problem is to recompute the bouquet from scratch, but most of the processing may turn out to be redundant. Therefore, developing incremental bouquet maintenance strategies is an interesting future research challenge.

Secondly, the bouquet identification overheads increase exponentially with dimensionality. Apart from the obvious amortization over repeated query invocations, we also described some mechanisms for reducing these overheads in Section 6.1. Further, a complex query does not necessarily imply a commensurately large number of error dimensions because: (i) The selectivities of base relation predicates of the form “*column op constant*” can be estimated accurately with current techniques; (ii) The join-selectivities for PK-FK joins can be estimated accurately if the entire PK-relation participates in the join; (iii) The partial derivatives of the POSP plan cost functions along each dimension can be computed on a low resolution mapping of the ESS, and any dimension with a small derivative across all the plans can be eliminated since its cost impact is marginal.

Thirdly, the identification of ESS dimensions may not always be straightforward. For example, in cyclic queries, different plans may combine predicates in different ways. One option to handle this scenario is to first construct the ESS using individual predicates as dimensions. Then, assuming that predicate independence holds, the selectivity of any predicate combination could be *inferred* using the existing values for the individual constituent predicates.

Given the above discussion, the bouquet approach is currently recommended specifically for providing response-time robustness in large archival read-only databases supporting complex decision-

support applications that are likely to suffer significant estimation errors. We expect that many of today’s OLAP installations may fall into this category.

In closing, we wish to highlight that the bouquet approach provides novel performance guarantees that open up new possibilities for robust query processing.

Acknowledgments. We thank the anonymous reviewers and S. Sudarshan, Prasad Deshpande, Srinivas Karthik, Sumit Neelam and Bruhathi Sundarmurthy for their constructive comments on this work.

9. REFERENCES

- [1] A. Aboulnaga and S. Chaudhuri, “Self-tuning Histograms: Building Histograms without Looking at Data”, *ACM SIGMOD Conf.*, 1999.
- [2] R. Avnur and J. Hellerstein, “Eddies: Continuously Adaptive Query Processing”, *ACM SIGMOD Conf.*, 2000.
- [3] B. Babcock and S. Chaudhuri, “Towards a Robust Query Optimizer: A Principled and Practical Approach”, *ACM SIGMOD Conf.*, 2005.
- [4] S. Babu, P. Bizarro and D. DeWitt, “Proactive Re-Optimization”, *ACM SIGMOD Conf.*, 2005.
- [5] P. Bizarro, N. Bruno, D. DeWitt, “Progressive Parametric Query Optimization”, *IEEE TKDE*, 21(4), 2009.
- [6] S. Chaudhuri, H. Lee and V. Narasayya, “Variance aware optimization of parameterized queries”, *ACM SIGMOD Conf.*, 2010.
- [7] S. Chaudhuri, V. Narasayya and R. Ramamurthy, “A Pay-As-You-Go Framework for Query Execution Feedback”, *PVLDB*, 1(1), 2008.
- [8] M. Chrobak, C. Kenyon, J. Noga and N. Young, “Incremental Medians via Online Bidding”, *Algorithmica*, 50(4), 2008.
- [9] F. Chu, J. Halpern and J. Gehrke, “Least Expected Cost Query Optimization: What Can We Expect”, *ACM PODS Conf.*, 2002.
- [10] R. Cole and G. Graefe, “Optimization of Dynamic Query Evaluation Plans”, *ACM SIGMOD Conf.*, 1994.
- [11] A. Deshpande, Z. Ives and V. Raman, “Adaptive Query Processing”, *Foundations and Trends in Databases*, Now Publishers, 2007.
- [12] A. Dutt and J. Haritsa, “Query Processing without Selectivity Estimation”, *Tech. Report TR-2014-01*, DSL/SERC, IISc, 2014, dsl.serc.iisc.ernet.in/publications/report/TR/TR-2014-01.pdf
- [13] G. Graefe et al, “Robust Query Processing (Dagstuhl Seminar 12321)”, *Dagstuhl Reports*, 2(8), 2012.
- [14] Harish D., P. Darera and J. Haritsa, “Identifying Robust Plans through Plan Diagram Reduction”, *PVLDB*, 1(1), 2008.
- [15] Harish D., P. Darera and J. Haritsa, “On the Production of Anorexic Plan Diagrams”, *VLDB Conf.*, 2007.
- [16] Y. Ioannidis and S. Christodoulakis, “On the Propagation of Errors in the Size of Join Results”, *ACM SIGMOD Conf.* 1991.
- [17] N. Kabra and D. DeWitt, “Efficient Mid-Query Re-Optimization of Sub-Optimal Query Execution Plans”, *ACM SIGMOD Conf.* 1998.
- [18] V. Markl et al, “Robust Query Processing through Progressive Optimization”, *ACM SIGMOD Conf.*, 2004.
- [19] G. Moerkotte, T. Neumann and G. Steidl, “Preventing Bad Plans by Bounding the Impact of Cardinality Estimation Errors”, *PVLDB*, 2(1), 2009.
- [20] T. Neumann and C. Galindo-Legaria, “Taking the Edge off Cardinality Estimation Errors using Incremental Execution”, *BTW Conf.*, 2013.
- [21] N. Polyzotis, “Selectivity-based partitioning: A Divide and Union Paradigm for Effective Query Optimization”, *ACM CIKM Conf.*, 2005.
- [22] P. Selinger et al, “Access Path Selection in a Relational Database Management System”, *ACM SIGMOD Conf.*, 1979.
- [23] M. Stillger, G. Lohman, V. Markl and M. Kandil, “LEO – DB2’s LEarning Optimizer”, *VLDB Conf.*, 2001.
- [24] W. Wu et al, “Predicting Query Execution Times: Are Optimizer Cost Models Really Usable?”, *IEEE ICDE Conf.*, 2013.
- [25] technet.microsoft.com/en-us/library/ms186954(v=sql.105).aspx
- [26] www.ibm.com/developerworks/data/library/tips/dm-0312yip/
- [27] www.oracle.com/ocom/groups/public/@otn/documents/webcontent/1963236.pdf
- [28] www.postgresql.org/docs/8.4/static/release.html
- [29] doxygen.postgresql.org/structInstrumentation.html